

Multivalued Dependencies

Winter 2006-2007

Lecture 22

Multivalued Attributes


- E-R schemas can have multivalued attributes
- 1NF requires only atomic attributes
 - Not a problem; translating to relational model leaves everything atomic

- Employee example:

employee(emp_id, emp_name)

emp_deps(emp_id, dependent)

emp_nums(emp_id, phone_num)



- What are the requirements on these schemas for what tuples must appear?

Multivalued Attributes (2)

- Example data:

<i>emp_id</i>	<i>emp_name</i>
125623	Rick

employee

<i>emp_id</i>	<i>dependent</i>
125623	Jeff
125623	Alice

emp_deps

<i>emp_id</i>	<i>phone_num</i>
125623	555-8888
125623	555-2222

emp_nums

- Every distinct value of multivalued attribute requires a separate tuple, including associated value of *emp_id*
- A consequence of 1NF, in fact!
 - If attributes could be nonatomic, could just store list of values in the appropriate column!
 - 1NF requires extra tuples to represent multivalues

Independent Multivalued Attributes

- Question is trickier when a schema stores several *independent* multivalued attributes
- Proposed combined schema:
 $employee(\underline{emp_id}, emp_name)$
 $emp_info(emp_id, dependent, phone_num)$
- What tuples must appear in emp_info ?
 - emp_info is a relation
 - If an employee has M dependents and N phone numbers, emp_info must contain $M \times N$ tuples
 - Every combination of the employee's dependents and their phone numbers

Independent Multivalued Attributes

- Example data:

<i>emp_id</i>	<i>emp_name</i>
125623	Rick

employee

<i>emp_id</i>	<i>dependent</i>	<i>phone_num</i>
125623	Jeff	555-8888
125623	Jeff	555-2222
125623	Alice	555-8888
125623	Alice	555-2222

emp_info

- Clearly has unnecessary redundancy
- Can't formulate functional dependencies to represent multivalued attributes
- Can't use BCNF or 3NF decompositions to eliminate redundancy in these cases

Dependencies

- Functional dependencies rule out what tuples can appear in a relation
 - If $A \rightarrow B$ holds, then tuples cannot have same value for A but different values for B
 - Also called equality-generating dependencies
- Multivalued dependencies specify what tuples must be present
 - To represent a multivalued attribute's values, a certain set of tuples *must* be present
 - Also called tuple-generating dependencies

Multivalued Dependencies

- Given a relation schema R
 - Attribute-sets $\alpha \in R, \beta \in R$
 - $\alpha \twoheadrightarrow \beta$ is a multivalued dependency
 - “ α multidetermines β ”
- Multivalued dependency $\alpha \twoheadrightarrow \beta$ holds on R if, in any legal relation $r(R)$:
 - For all pairs of tuples t_1 and t_2 in r such that $t_1[\alpha] = t_2[\alpha]$
 - There also exists tuples t_3 and t_4 in r such that:
 - $t_1[\alpha] = t_2[\alpha] = t_3[\alpha] = t_4[\alpha]$
 - $t_1[\beta] = t_3[\beta]$ and $t_2[\beta] = t_4[\beta]$
 - $t_1[R - \beta] = t_4[R - \beta]$ and $t_2[R - \beta] = t_3[R - \beta]$

Multivalued Dependencies (2)

- Multivalued dependency $\alpha \twoheadrightarrow \beta$ holds on R if, in any legal relation $r(R)$:
 - For all pairs of tuples t_1 and t_2 in r such that $t_1[\alpha] = t_2[\alpha]$
 - There also exists tuples t_3 and t_4 in r such that:
 - $t_1[\alpha] = t_2[\alpha] = t_3[\alpha] = t_4[\alpha]$
 - $t_1[\beta] = t_3[\beta]$ and $t_2[\beta] = t_4[\beta]$
 - $t_1[R - \beta] = t_4[R - \beta]$ and $t_2[R - \beta] = t_3[R - \beta]$

- Pictorially:

	α	β	$R - (\alpha \cup \beta)$
t_1	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$a_{j+1} \dots a_n$
t_2	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$b_{j+1} \dots b_n$
t_3	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$b_{j+1} \dots b_n$
t_4	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$a_{j+1} \dots a_n$

Multivalued Dependencies (3)

- Multivalued dependency:

	α	β	$R - (\alpha \cup \beta)$
t_1	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$a_{j+1} \dots a_n$
t_2	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$b_{j+1} \dots b_n$
t_3	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$b_{j+1} \dots b_n$
t_4	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$a_{j+1} \dots a_n$

- If $\alpha \twoheadrightarrow \beta$ then $R - (\alpha \cup \beta)$ is independent of this
 - Every distinct value of β must be associated once with every distinct value of $R - (\alpha \cup \beta)$
- Let $\gamma = R - (\alpha \cup \beta)$
 - If $\alpha \twoheadrightarrow \beta$ then also $\alpha \twoheadrightarrow \gamma$
 - $\alpha \twoheadrightarrow \beta$ implies $\alpha \twoheadrightarrow \gamma$
 - Sometimes written $\alpha \twoheadrightarrow \beta \mid \gamma$

Trivial Multivalued Dependencies

- $\alpha \twoheadrightarrow \beta$ is a trivial multivalued dependency on R if all relations $r(R)$ satisfy the dependency
- Specifically, $\alpha \twoheadrightarrow \beta$ is trivial if $\beta \subseteq \alpha$, or if $\alpha \cup \beta = R$
- Employee examples:
 - For schema $emp_deps(emp_id, dependent)$, $emp_id \twoheadrightarrow dependent$ is trivial
 - For $emp_info(emp_id, dependent, phone_num)$, $emp_id \twoheadrightarrow dependent$ is not trivial

Inference Rules

- Can reason about multivalued dependencies, just like functional dependencies
- Example inference rules:
 - Complementation rule:
 - If $\alpha \twoheadrightarrow \beta$ holds on R , then $\alpha \twoheadrightarrow R - (\alpha \cup \beta)$ holds
 - Multivalued augmentation rule:
 - If $\alpha \twoheadrightarrow \beta$ holds, and $\gamma \supseteq \delta$, then $\alpha\gamma \twoheadrightarrow \beta\delta$
 - Multivalued transitivity rule:
 - If $\alpha \twoheadrightarrow \beta$ and $\beta \twoheadrightarrow \gamma$ holds, then $\alpha \twoheadrightarrow \gamma - \beta$ holds
 - Coalescence rule:
 - If $\alpha \twoheadrightarrow \beta$ and there exists γ such that $\gamma \cap \beta = \emptyset$, and $\gamma \rightarrow \delta$, and $\beta \supseteq \delta$, then $\alpha \rightarrow \delta$

Functional Dependencies

- Functional dependencies are also multivalued dependencies
 - If $\alpha \rightarrow \beta$, then $\alpha \twoheadrightarrow \beta$ too
 - Caveat: each value of α has at most one associated value for β
- Usually, functional dependencies are not stated as multivalued dependencies
 - Because of additional caveat; not obvious in notation
 - Also because functional dependencies are just easier to reason about!

Closures and Restrictions

- For a set D of functional and multivalued dependencies, can compute closure D^+
 - Use inference rules for both functional and multivalued dependencies to compute closure
- Sometimes need the restriction of D^+ to a relation schema R , too
- The restriction of D to a schema R includes:
 - All functional dependencies in D^+ that include only attributes in R
 - All multivalued dependencies of the form $\alpha \twoheadrightarrow \beta \cap R$, where $\alpha \subseteq R$, and $\alpha \twoheadrightarrow \beta$ is in D^+

Fourth Normal Form

- Given:
 - Relation schema R
 - Set of functional and multivalued dependencies D
- R is in 4NF with respect to D if:
 - For all multivalued dependencies $\alpha \twoheadrightarrow \beta$ in D^+ , where $\alpha \in R$ and $\beta \in R$, at least one of the following holds:
 - $\alpha \twoheadrightarrow \beta$ is a trivial multivalued dependency
 - α is a superkey for R
 - Note: If $\alpha \rightarrow \beta$ then $\alpha \twoheadrightarrow \beta$
- A database design is in 4NF if all schemas in the design are in 4NF

4NF and BCNF

- Main difference between 4NF and BCNF is use of multivalued dependencies instead of functional dependencies
- Every schema in 4NF is also in BCNF
 - If a schema is not in BCNF then there is a nontrivial functional dependency $\alpha \rightarrow \beta$ such that α is not a superkey for R
 - If $\alpha \rightarrow \beta$ then $\alpha \twoheadrightarrow \beta$

4NF Decompositions

- Decomposition rule very similar to BCNF
- If schema R is not in 4NF with respect to a set of multivalued dependencies D :
 - There is some nontrivial dependency $\alpha \twoheadrightarrow \beta$ in D^+ where $\alpha \subseteq R$ and $\beta \subseteq R$, and α is not a superkey of R
 - Also constrain that $\alpha \cap \beta = \emptyset$
 - Replace R with two new schemas:
 - $R_1 = (\alpha \cup \beta)$
 - $R_2 = (R - \beta)$

Employee Information Example

- Combined schema:
 $employee(\underline{emp_id}, emp_name)$
 $emp_info(emp_id, dependent, phone_num)$
 - Also have these dependencies:
 - $emp_id \rightarrow emp_name$
 - $emp_id \twoheadrightarrow dependent$
 - $emp_id \twoheadrightarrow phone_num$
- emp_info is not in 4NF
 - Following rules for 4NF decomposition produces:
 $(emp_id, dependent)$
 $(emp_id, phone_num)$

Lossless Decompositions

- Can also define lossless decomposition with multivalued dependencies
 - R_1 and R_2 form a lossless decomposition of R if at least one of these dependencies is in D^+ :

$$R_1 \cap R_2 \twoheadrightarrow R_1$$

$$R_1 \cap R_2 \twoheadrightarrow R_2$$

Beyond Fourth Normal Form?

- Several other normal forms with various constraints
- Some define new dependencies, such as:
 - Join dependencies: a more general form of multivalued dependencies
 - Inclusion dependencies: for representing foreign key constraints
- Fifth normal form (5NF) includes join dependencies
 - Also known as project-join normal form (PJNF)
- Domain-key normal form (DKNF) is an even more general normal form
 - Takes domain constraints into account
 - Can state other normal forms as special cases of DKNF
- Higher normal forms are much harder to reason about
 - Not widely used, although often good goals to aim for!

Normalized Schemas

- Normalized schemas have certain features
 - Usually more relations, to eliminate redundancy
 - Usually need to join several relations together to retrieve a particular result
 - Manipulating the data is usually easier, because there is very little redundancy
- Sometimes a database application needs high performance query support
 - Normalized database's layout doesn't facilitate fast query operations

Faster Query Evaluation

- Sometimes database designers will denormalize a database schema
 - Intentionally design a schema to violate a normal form, in order to be faster
 - Usually requires more development effort, to keep redundant data synchronized
 - For systems requiring high transaction throughput, may be the only option!
 - Should usually consider this as a last resort
 - Try other options first!

Materialized Views

- Materialized views can provide a denormalized view of a normalized schema
 - Create a normalized schema
 - Define widely used views against schema
 - The database stores the view's results on disk
- Database itself keeps view in sync with underlying tables
 - Unlike a denormalized schema, developers don't have to worry about managing things
- Not all databases provide materialized views

Materialized Views (2)

- Materialized views can become expensive if data changes frequently
 - Especially when a value in the underlying relation is duplicated multiple times in the view
- One solution: update view periodically
 - e.g. view is recomputed and stored hourly or daily
- For applications where underlying data changes very frequently, a denormalized schema may be faster

Unique-Role Assumption

- SQL requires every table has unique column names
- Column names indicate what values represent
 - Different tables can have columns with the same name, but different meanings
- The unique-role assumption:
 - Every attribute name in a schema has a unique meaning in the database
 - Some obvious benefits:
 - Clearly indicates foreign key relationships between tables
 - No ambiguity in what an attribute name means in a query, or in its result
 - Can use **NATURAL JOIN** syntax for most queries

Unique-Role Assumption (2)

- A common example:

assignment(id, shortname, url, perfectscore, due)

submission(id, assignment, graded, score)

– For each table, *id* is the unique identifier for that table's rows

– Foreign key relationships become awkward

- *submission.assignment* is foreign key to *assignment.id*

- A better design:

assignment(asn_id, shortname, url, perfectscore, due)

submission(sub_id, asn_id, graded, score)

– Follows unique-role assumption guideline

Unique-Role Assumption (3)

- One situation where unique-role assumption is harder to follow:
 - A relation with a foreign key reference to itself
- Example:
 - employee(emp_id, emp_name, salary)*
 - manages(emp_id, manager_id)*
 - *manager_id* has same domain as *emp_id*
 - Simply can't give both attributes the same name!
(Nor would you necessarily want to...)
- Usually occurs infrequently in schema designs
 - Only certain circumstances where relations have foreign key references to themselves

General Naming Concerns

- Unique-role assumption is one specific aspect of a more general concern
- Database schemas contain lots of names!
 - Table names, column names
 - Possibly constraint names and stored procedure names, as well
 - Names become hard to change after a while
- Database schemas should follow a unified naming convention
 - Singular vs. plural table names
 - e.g. *employee* vs. *employees*
 - Unique-role assumption for column names, etc.

Review

- Covered normal forms for database schemas
 - Patterns to follow that make for “good” schemas
- Goals:
 - Representation must be complete
 - Eliminate unnecessary redundancy
- Formalized the concept using dependencies
 - Functional dependencies, multivalued dependencies
 - Rules of inference for dependencies
 - In some cases, algorithms for generating normalized schemas

Review, and Preview!

- Designing normalized schemas can involve trade-offs
 - Can mean slower performance, but easier development and maintenance
 - Some ways to work around slow performance
 - Sometimes, only choice is a denormalized schema
- Next time:
 - How to maximize performance of query evaluation?
 - How do databases actually evaluate queries?
 - What tools do databases provide to improve performance?